



SAS Industry Orientation Program
Department of Statistics
University of Lucknow



SAS Programming Course

(March 16, 2024 – May 04, 2024)

Department of Statistics
University of Lucknow

Disclaimer:



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

The topics, thoughts, case studies etc. in the discussions are sole belongs to me and do not belong or represent organization where we work.

Agenda for today:

- Introduction to Machine Learning
- Frequent terms used in Machine Learning World
- Use cases of Machine Learning



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Building Data Science Solutions

Introduction to Data Science terminologies

1. Data Science
2. Algorithm
3. Features/Variables
4. Artificial Intelligence
5. Machine Learning
6. Deep Learning
7. Natural Language Processing
8. Prediction
9. Regression
10. Classification
11. Supervised Machine Learning
12. Un-Supervised Machine Learning
13. Clustering
14. Neural Networks
15. Overfitting/Underfitting
16. Encoding
17. Embeddings
18. Sentiment Analysis
19. Large Language Models
20. A/B Testing



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Building Data Science Solutions

Data Science



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Data science for business involves using **data analysis techniques** and **technologies** to help businesses make **better decisions**, **solve problems**, and achieve their goals. It involves **collecting, processing, analyzing, and interpreting large volumes of data** to uncover **insights** that can drive **strategic and operational decisions**. These insights can range from understanding **customer behavior** and preferences to **optimizing operations, identifying market trends, predicting future outcomes**, and much more. Ultimately, data science for business aims to **leverage data as a valuable asset** to **improve efficiency, effectiveness, and competitiveness** in the marketplace.

Building Data Science Solutions

Algorithms

An algorithm is a **step-by-step set of instructions or a procedure** designed to **solve a particular problem** or **perform a specific task**.

In the context of business and management, algorithms are often used in various software applications and systems to **automate processes, analyze data, and make decisions**. They can range from simple rules-based procedures to complex mathematical models, all aimed at **improving efficiency, optimizing resources, and achieving organizational objectives**. Essentially, algorithms help streamline operations, enhance decision-making processes, and drive innovation within a company.



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Building Data Science Solutions

Features

In machine learning, features (also known as input variables or independent variables) are the individual measurable properties or characteristics of the data that are used to make predictions or decisions.

Features are the inputs to a machine learning model and play a crucial role in determining its performance.

Features are fundamental components of machine learning models, and careful selection, preprocessing, and engineering of features are essential for building accurate and effective predictive models.



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Building Data Science Solutions

Artificial Intelligence (AI)


From a management perspective, Artificial Intelligence (AI) can be defined as the integration of advanced computing technologies that enable machines to perform tasks that typically require human intelligence. This includes tasks such as understanding natural language, recognizing patterns in data, making decisions, and even learning from experience.

AI represents a transformative force that can drive business growth, improve decision-making, and enhance competitiveness in today's rapidly evolving digital landscape.

Effective management of AI initiatives involves leveraging its capabilities to achieve strategic objectives while addressing ethical and societal considerations.



SAS Industry Orientation Program
Department of Statistics
University of Lucknow



Building Data Science Solutions


Machine Learning (ML)

Machine learning can be defined as a **powerful technology** that enables computers to learn from **data and make predictions or decisions** without being **explicitly programmed** to do so. In simpler terms, it's like having a smart assistant that can **analyze large amounts of data to identify patterns, trends, and insights** that can help businesses make better **decisions, improve processes, and achieve their goals**.

Machine learning empowers businesses to leverage their **data assets** more effectively, driving **innovation, efficiency, and competitiveness** in today's data-driven economy. By harnessing the power of machine learning, businesses can unlock **new opportunities, solve complex problems, and achieve their strategic objectives**.




SAS Industry Orientation Program
Department of Statistics
University of Lucknow



Building Data Science Solutions

Deep Learning (DL)

Deep learning can be defined as **an advanced subset of machine learning** that mimics the workings of the **human brain** to process and understand **complex data**. Deep learning algorithms are designed to **automatically learn representations of data** through multiple layers of **abstraction**, allowing them to **extract intricate patterns** and make highly accurate predictions.



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Building Data Science Solutions

AI, ML and DL in a nutshell



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Aspect	Artificial Intelligence (AI)	Machine Learning (ML)	Deep Learning (DL)
Definition	A broad field of computer science that aims to create intelligent machines capable of performing tasks that typically require human intelligence.	A subset of AI that involves developing algorithms that allow computers to learn from data and make predictions or decisions without being explicitly programmed.	A subset of machine learning that uses neural networks with multiple layers (deep architectures) to learn complex patterns and representations from data.
Approach	May involve various techniques including rule-based systems, expert systems, and symbolic reasoning.	Utilizes algorithms that learn from data and iteratively improve performance over time without being explicitly programmed.	Utilizes artificial neural networks with multiple layers to automatically learn representations of data through abstraction.
Scope	Encompasses a wide range of applications including natural language processing, robotics, computer vision, and more.	Widely used in applications such as predictive analytics, pattern recognition, classification, and clustering.	Primarily used in tasks such as image recognition, speech recognition, natural language processing, and more, where large amounts of data are available.
Model Complexity	Can range from simple rule-based systems to complex cognitive architectures.	Can range from simple linear regression models to complex ensemble methods and deep neural networks.	Typically involves complex neural network architectures with multiple hidden layers, requiring significant computational resources.

Building Data Science Solutions

AI, ML and DL in a nutshell



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Aspect	Artificial Intelligence (AI)	Machine Learning (ML)	Deep Learning (DL)
Data Requirements	May require structured and unstructured data from various sources depending on the application.	Requires labeled or unlabeled training data for algorithm training and validation.	Requires large amounts of labeled or unlabeled data for training, which can be a limiting factor in some applications.
Training Process	May involve manual programming, knowledge engineering, or learning from experience.	Involves training algorithms on historical data to learn patterns and make predictions or decisions.	Involves training deep neural networks on large datasets using techniques such as backpropagation and stochastic gradient descent.
Interpretability	Outputs may not always be explainable or interpretable by humans.	Outputs can often be interpreted and understood based on the model's underlying logic or decision rules.	Outputs may be less interpretable due to the complex, hierarchical nature of deep neural networks.
Applications	Common applications include virtual assistants, autonomous vehicles, recommendation systems, and more.	Widely used in industries such as finance, healthcare, retail, and marketing for tasks such as fraud detection, customer segmentation, and demand forecasting.	Common applications include image and speech recognition, natural language processing, autonomous driving, and more.
Examples	IBM Watson, Siri, self-driving cars, facial recognition systems.	Logistic regression, decision trees, random forests, support vector machines.	Convolutional neural networks (CNNs), recurrent neural networks (RNNs), generative adversarial networks (GANs).

Building Data Science Solutions


Natural Language Processing (NLP)



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Natural Language Processing (NLP) is a technology that enables computers to **understand, interpret, and generate human language** in a way that's similar to how people communicate.

Imagine you have a lot of **text-based information**, like **emails, social media posts, or customer reviews**. NLP helps businesses make sense of all this text by allowing computers to understand the meaning behind the words. NLP helps businesses unlock **valuable insights from text-based data**, improve **customer interactions**, **automate tasks**, and expand their global reach.



Building Data Science Solutions

Supervised & Unsupervised Machine Learning




SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Supervised machine learning is a type of **machine learning** where the **algorithm learns from labeled data**, meaning the **input data** is paired with the **corresponding correct output**.

The algorithm learns to make predictions or decisions based on this labeled dataset.

Unsupervised machine learning is a type of machine learning where the algorithm learns from **unlabeled data**, meaning the **input data does not have corresponding output labels**.

The goal of unsupervised learning is to find **patterns, relationships, or structures** within the data without explicit guidance.



Building Data Science Solutions

Supervised & Unsupervised Machine Learning in nutshell



AS Industry Orientation Program
Department of Statistics
University of Lucknow

Aspect	Supervised Machine Learning	Unsupervised Machine Learning
Input Data	Labeled data, consisting of input features and corresponding output labels.	Unlabeled data, consisting only of input features without corresponding output labels.
Learning Approach	Learns from labeled examples to predict or classify new instances.	Learns from unlabeled data to discover patterns, relationships, or structures within the data.
Goal	To make predictions or decisions based on input features.	To find patterns, groupings, or anomalies within the data.
Examples	Classification, Regression.	Clustering, Dimensionality Reduction, Anomaly Detection.
Evaluation Metrics	Accuracy, Precision, Recall (for classification).	No clear output labels, evaluation often subjective or based on domain knowledge.
Model Complexity	Typically more complex models with higher capacity to learn intricate patterns.	Simpler models, often relying on clustering or dimensionality reduction techniques.
Feedback Loop	Requires labeled data for training, feedback loop involves comparing predictions to actual labels.	Does not require labeled data, feedback loop may involve adjusting model parameters based on unsupervised learning techniques.
Use Cases	Predictive modeling, Sentiment Analysis, Image Recognition.	Customer Segmentation, Anomaly Detection, Data Exploration.
Interpretability	Often more interpretable, as predictions are based on labeled examples and can be traced back to input features.	May be less interpretable, as patterns and groupings are discovered without explicit guidance.
Data Requirements	Requires labeled data, which can be expensive and time-consuming to obtain.	Can work with unlabeled data, potentially reducing data collection and labeling efforts.

Building Data Science Solutions

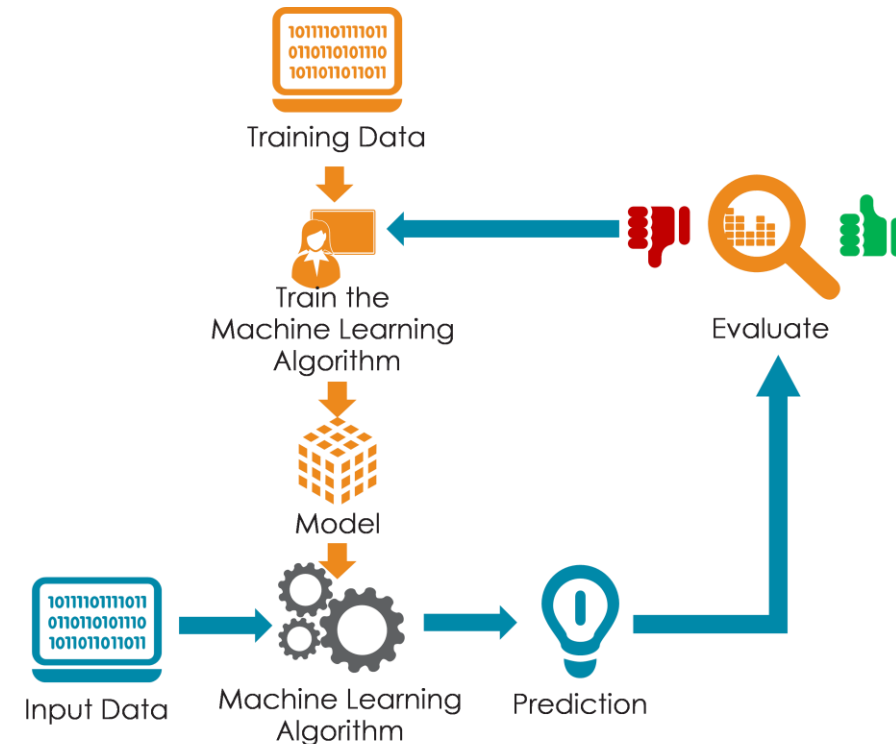
Prediction



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Imagine you're a business owner of an online retail store. You have a lot of **historical data** about your **customers**, such as their **purchase history**, **demographics**, and **website behavior**. Machine learning prediction can help you **forecast/predict** what your customers are likely to do in the future based on this data.

Machine learning prediction empowers businesses to **anticipate future trends, behaviors**, and **outcomes** based on historical data. By leveraging these predictions, businesses can make more **informed decisions**, **improve customer experiences**, and **drive growth and profitability**.



Building Data Science Solutions


Regression



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Regression problems are a type of supervised learning problem in machine learning where the goal is to predict a continuous numerical value based on input variables. In other words, regression models are used when the target or dependent variable is a real-valued number.

Regression problems involve predicting a continuous numerical value based on input variables, and regression analysis is widely used in various fields for prediction and forecasting tasks.



Building Data Science Solutions


Classification



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Classification problems are a type of **supervised learning problem** in machine learning where the goal is to **predict the category or class label** of a new observation based on input variables. In other words, classification models are used when the **target or dependent variable is categorical**.

Classification problems involve predicting the category or class label of a new observation based on input variables, and classification models are widely used in various fields for tasks such as **sentiment analysis, image recognition, and fraud detection**.



Building Data Science Solutions

Clustering



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Clustering in machine learning is like **organizing** your business's customer base into different **groups based on similarities**. It's a way to group together customers who share **common characteristics or behaviors**, even if you didn't know about these groups beforehand.

Clustering in machine learning is a powerful tool for businesses to **understand their customers** and **tailor their strategies** accordingly. It helps you uncover **hidden patterns**, **group customers into meaningful segments**, and make data-driven decisions that drive growth and profitability.



Building Data Science Solutions


Artificial Neural networks (ANN)



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Neural networks are like **virtual brains for computers**. Just as our brains have **interconnected neurons** that process information, neural networks are made up of **artificial neurons** called **nodes**, organized into layers.

Neural networks are computer systems inspired by the human brain. They learn from **data, recognize patterns, and make predictions**, making them valuable tools for solving complex problems and driving innovation in various fields.



Building Data Science Solutions

Supervised Machine Learning use cases



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Algorithm	Description	Use Cases in Business
Linear Regression	Predicts a continuous numerical value based on input features by fitting a linear equation to the data.	Sales forecasting, price optimization, demand prediction.
Logistic Regression	Predicts the probability of a binary outcome based on input features using a logistic function.	Customer churn prediction, credit risk assessment, fraud detection.
Decision Trees	Builds a tree-like structure to classify or regress data by recursively splitting based on feature values.	Customer segmentation, product recommendation, risk analysis.
Random Forest	Ensemble learning method that combines multiple decision trees to improve prediction accuracy.	Marketing campaign targeting, credit scoring, anomaly detection.
Support Vector Machines	Finds the optimal hyperplane that best separates data points into different classes in a high-dimensional space.	Image classification, sentiment analysis, text categorization.
Naive Bayes	Probabilistic classifier based on Bayes' theorem that assumes independence between features.	Email spam detection, document classification, sentiment analysis.
K-Nearest Neighbors (KNN)	Classifies data points based on the majority class among their k nearest neighbors in feature space.	Customer segmentation, recommendation systems, anomaly detection.
Neural Networks	Deep learning models composed of interconnected layers of nodes that learn complex patterns from data.	Image recognition, speech recognition, natural language processing.

Building Data Science Solutions

Un-Supervised Machine Learning use cases



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Algorithm	Description	Use Cases in Business
K-Means Clustering	Divides data into 'k' clusters based on similarity, with each cluster represented by its centroid.	Customer segmentation, Market basket analysis, Anomaly detection, Product recommendation.
Hierarchical Clustering	Builds a hierarchy of clusters by either merging or splitting existing clusters based on similarity.	Customer segmentation, Taxonomy creation, Fraud detection, Image segmentation.
DBSCAN	Groups together data points that are closely packed based on density, identifying noise as outliers.	Anomaly detection, Customer segmentation, Geographic clustering, Fraud detection.
Principal Component Analysis (PCA)	Reduces the dimensionality of the data while preserving most of its variance, by transforming it into a new set of orthogonal variables.	Dimensionality reduction, Feature extraction, Data visualization, Anomaly detection.
t-Distributed Stochastic Neighbor Embedding (t-SNE)	Reduces high-dimensional data to a lower-dimensional space while preserving local structure.	Data visualization, Clustering validation, Exploratory data analysis, Feature engineering.
Gaussian Mixture Models (GMM)	Represents a mixture of several Gaussian distributions to model complex data distributions.	Clustering, Density estimation, Anomaly detection, Image segmentation.

Building Data Science Solutions

Artificial Neural Network (ANN) use cases



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Neural Network Architecture	Description	Use Cases
Feedforward Neural Network (FNN)	Consists of input, hidden, and output layers where information flows in one direction from input to output. Each neuron in a layer is connected to all neurons in the subsequent layer.	Classification, Regression, Pattern recognition, Credit scoring.
Convolutional Neural Network (CNN)	Designed for processing grid-like data such as images. Utilizes convolutional layers to extract features hierarchically and pooling layers to reduce dimensionality.	Image classification, Object detection, Facial recognition, Medical image analysis.
Recurrent Neural Network (RNN)	Contains recurrent connections allowing information to persist over time. Suitable for sequence data processing where the output depends on the previous inputs.	Natural language processing, Speech recognition, Time series prediction, Sentiment analysis.
Long Short-Term Memory (LSTM)	A type of RNN architecture with memory cells and gating mechanisms to capture long-range dependencies and mitigate vanishing gradient problem.	Text generation, Machine translation, Stock price prediction, Video analysis.
Gated Recurrent Unit (GRU)	Similar to LSTM but with a simpler architecture. It combines the forget and input gates into a single update gate.	Speech recognition, Language modeling, Gesture recognition, Handwriting recognition.
Autoencoder	Consists of an encoder and a decoder. Encoder compresses input data into a latent representation, and decoder reconstructs the original input from the latent representation.	Dimensionality reduction, Feature learning, Anomaly detection, Denoising data.
Generative Adversarial Network (GAN)	Comprises two neural networks, a generator and a discriminator, trained simultaneously in a competitive manner. Generator creates realistic data samples, and discriminator distinguishes between real and fake samples.	Image generation, Data augmentation, Style transfer, Super-resolution.

Building Data Science Solutions

Overfitting and Underfitting



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

In machine learning, **overfitting** happens when a **model learns too much** from the **training data** and **captures noise or random fluctuations** in the data as if they were real patterns.

As a result, the **model performs very well on the training data** but **fails to generalize to new, unseen data**.

In machine learning, **underfitting** happens when a model is **too simple to capture the underlying structure of the data**. The **model fails to learn from the training data** and **performs poorly both on the training data and on new, unseen data**.

Building Data Science Solutions

Encoding



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

In machine learning, "encodings" refer to the process of representing **categorical data** (data that consists of categories or labels) in a **numerical format** that machine learning algorithms can **understand and process effectively**. Encodings are necessary because most **machine learning algorithms require numerical input data**.

These encoding techniques allow machine learning algorithms to **effectively process categorical data, enabling them** to learn from and make predictions on datasets that include **categorical variables**. Choosing the **appropriate** encoding technique depends on the **nature of the categorical variable** and the specific requirements of the machine learning task.

Building Data Science Solutions

Encoding - Examples



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Encoding Method	Description	Use Cases
Label Encoding	Assigns a unique numerical label to each category in a categorical variable.	Ordinal categorical variables, such as "Low," "Medium," "High" levels of a feature.
One-Hot Encoding	Creates binary dummy variables for each category, where each category is represented by a binary vector.	Nominal categorical variables, such as colors ("Red," "Blue," "Green").
Binary Encoding	Represents each category as a binary number, with each digit of the binary number becoming a feature.	Large categorical variables, where one-hot encoding would lead to high dimensionality.
Target Encoding	Replaces each category with the mean (or other statistic) of the target variable for that category.	Variables with high cardinality, such as geographical regions or product categories.
Frequency Encoding	Replaces each category with the frequency (count) of occurrences of that category in the dataset.	Rare event detection, anomaly detection, or when the frequency of occurrence is relevant information.
Ordinal Encoding	Similar to label encoding but assigns numerical labels based on the ordinality (order) of categories.	Ordinal categorical variables, where the order of categories is meaningful, such as "Low," "Medium," "High" levels of a feature.

Building Data Science Solutions

Word Embeddings



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

In simple terms, **word embeddings** in natural language processing (NLP) are **numerical representations** of words that capture their **meaning and context** in a way that computers can understand.

Imagine you're teaching a computer to understand language. Instead of treating each word as a standalone entity, word embeddings allow the computer to **understand relationships between words based** on their usage in context.

Word embeddings provide a way to represent words as **numerical vectors** that capture their **meaning and context in a high-dimensional space**. They are a fundamental component of many **NLP models, enabling computers** to understand and work with human language more intelligently.

Building Data Science Solutions

Words Embeddings - Examples



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Word Embedding Method	Description	Use Cases
Word2Vec	A shallow neural network model that learns word embeddings by predicting words in context.	Natural language processing tasks such as sentiment analysis, document classification, and machine translation.
GloVe (Global Vectors for Word Representation)	Uses co-occurrence statistics to learn word embeddings, emphasizing global context.	Similar use cases as Word2Vec, often preferred for tasks where semantic relationships between words are crucial, such as word analogy tasks.
FastText	An extension of Word2Vec that represents words as bags of character n-grams, allowing for better handling of out-of-vocabulary words and morphologically rich languages.	Text classification, Language modeling, Named entity recognition, Text similarity tasks.
ELMo (Embeddings from Language Models)	Generates word embeddings by combining deep bidirectional language model representations.	Natural language understanding tasks such as question answering, sentiment analysis, and named entity recognition.
BERT (Bidirectional Encoder Representations from Transformers)	Utilizes transformer architecture to pre-train a deep bidirectional language model on large text corpora.	Cutting-edge natural language processing tasks such as text generation, question answering, and language understanding.
GPT (Generative Pre-trained Transformer)	A transformer-based model trained on a diverse range of internet text to generate coherent and contextually relevant text.	Text generation, Conversational agents, Content creation, Language understanding.

Building Data Science Solutions

Large Language Models (LLMs)



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Large Language Models (LLMs) are powerful **artificial intelligence models** trained on **massive amounts of data**. These models, such as **OpenAI's GPT** (Generative Pre-trained Transformer) series or **Google's BERT** (Bidirectional Encoder Representations from Transformers), are capable of understanding and generating human-like text at scale.

Large Language Models offer **tremendous potential for businesses across industries** to enhance their **operations, improve customer experiences, drive innovation, and gain a competitive edge** in today's data-driven and digitally connected world. However, it's essential for businesses to consider ethical and privacy implications when deploying LLMs and ensure responsible and transparent use of AI technologies.

Building Data Science Solutions

Large Language Models (LLMs) – Use Cases



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Large Language Model	Developer	Use Cases
LaMDA	Google AI	<ul style="list-style-type: none">• Chatbots and virtual assistants• Dialogue generation• Summarization of factual topics
GPT-3	OpenAI	<ul style="list-style-type: none">• Creative text formats (poems, code, scripts)• Machine translation• Question answering• Content generation
Megatron-Turing NLG	NVIDIA	<ul style="list-style-type: none">• Scientific research paper generation• Summarization of complex information• Extractive question answering
WuDao 2.0	BAAI (China)	<ul style="list-style-type: none">• Machine translation (especially Chinese languages)• Dialogue generation• Question answering
Jurassic-1 Jumbo	AI21 Labs	<ul style="list-style-type: none">• Scientific writing assistance• Different creative text formats• Summarization• Question answering
LLaMA	Meta AI	<ul style="list-style-type: none">• Generating different creative text formats• Machine translation (with focus on cultural context)• Question answering

Building Data Science Solutions

A/B Testing



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

A/B testing, also known as **split testing**, is a method used to compare two versions of a **process, webpage or app** to determine which one performs better. It is commonly used in **marketing, user experience design, and product development** to optimize various aspects of a product or service.

- Website Design and Layout Optimization
- Email Marketing Campaigns
- Product Features and User Interface (UI) Enhancements
- Pricing and Promotions Optimization
- Mobile App Optimization



SAS Industry Orientation Program
Department of Statistics
University of Lucknow

Thank You.

Have a Question:

Anant Prakash Awasthi

M: +91-88846-52929

E: anant.awasthi@outlook.com

Dhananjay Srivastava

M: +91-80901-42364

E: ghananjay.srivastava@sscr.co.in